

Clothing Classification

Description

Fashion training set consists of 70,000 images divided into 60,000 training and 10,000 testing samples. Dataset sample consists of 28x28 grayscale image, associated with a label from 10 classes.

The 10 classes are as follows:

- 0 = T-shirt/top
- 1 = Trouser
- 2 = Pullover
- 3 = Dress
- 4 = Coat
- 5 = Sandal
- 6 = Shirt
- 7 = Sneaker
- 8 = Bag
- 9 = Ankle boot

Each image is 28 pixels in height and 28 pixels in width, for a total of 784 pixels in total. Each pixel has a single pixel-value associated with it, indicating the lightness or darkness of that pixel, with higher numbers meaning darker. This pixel-value is an integer between 0 and 255.

Importing the dataset

```
In [1]: # Importing required libraries

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [3]: # Let's create training and testing dataframes using CSV files

training_df = pd.read_csv('train_data.csv')
testing_df = pd.read_csv('test_data.csv')
```

Visualizing the data

```
In [4]: # Let's check first 5 rows of training dataframe

training_df.head()
```

Out[4]:

	label	pixel1	pixel2	pixel3	pixel4	pixel5	pixel6	pixel7	pixel8	pixel9	...	pixel775	pixel776	pi
0	2	0	0	0	0	0	0	0	0	0	...	0	0	
1	9	0	0	0	0	0	0	0	0	0	...	0	0	
2	6	0	0	0	0	0	0	0	5	0	...	0	0	
3	0	0	0	0	1	2	0	0	0	0	...	3	0	
4	3	0	0	0	0	0	0	0	0	0	...	0	0	

5 rows × 785 columns

```
In [5]: # Let's check last 5 rows of training dataframe

training_df.tail()
```

Out[5]:

	label	pixel1	pixel2	pixel3	pixel4	pixel5	pixel6	pixel7	pixel8	pixel9	...	pixel775	pixel776
59995	9	0	0	0	0	0	0	0	0	0	...	0	0
59996	1	0	0	0	0	0	0	0	0	0	...	73	0
59997	8	0	0	0	0	0	0	0	0	0	...	160	160
59998	8	0	0	0	0	0	0	0	0	0	...	0	0
59999	7	0	0	0	0	0	0	0	0	0	...	0	0

5 rows × 785 columns

Observations

- 1. There are 784 columns with various pixel values of different images
- 2. First column (i.e. label) is the column to be predicted. We will ignore the same initially.
- 3. There are 60000 observations in training dataset
- 4. All pixel values are between 0 to 255

```
In [6]: # Let's check first 5 rows of testing dataframe

testing_df.head()
```

Out[6]:

	label	pixel1	pixel2	pixel3	pixel4	pixel5	pixel6	pixel7	pixel8	pixel9	...	pixel775	pixel776	pixel777
0	0	0	0	0	0	0	0	0	9	8	...	103	87	0
1	1	0	0	0	0	0	0	0	0	0	...	34	0	0
2	2	0	0	0	0	0	0	14	53	99	...	0	0	0
3	2	0	0	0	0	0	0	0	0	0	...	137	126	0
4	3	0	0	0	0	0	0	0	0	0	...	0	0	0

5 rows × 785 columns

```
In [7]: # Let's check last 5 rows of testing dataframe

testing_df.tail()
```

Out[7]:

	label	pixel1	pixel2	pixel3	pixel4	pixel5	pixel6	pixel7	pixel8	pixel9	...	pixel775	pixel776	pixel777
9995	0	0	0	0	0	0	0	0	0	0	...	32	23	0
9996	6	0	0	0	0	0	0	0	0	0	...	0	0	0
9997	8	0	0	0	0	0	0	0	0	0	...	175	172	0
9998	8	0	1	3	0	0	0	0	0	0	...	0	0	0
9999	1	0	0	0	0	0	0	0	140	119	...	111	95	0

5 rows × 785 columns

Observations

- 1. There are 784 columns with various pixel values of different images
- 2. There are 10000 observations in testing dataset
- 3. All pixel values are between 0 to 255

```
In [8]: # Let's create numpy arrays for training and testing dataframes
# using float values for further processing
```

```
trainData = np.array(training_df, dtype='float32')
testData = np.array(testing_df, dtype='float32')
```

```
In [9]: # Let's check the training data and it's shape
```

```
print(trainData)
print('\n')
print(trainData.shape)
```

```
[[2. 0. 0. ... 0. 0. 0.]
 [9. 0. 0. ... 0. 0. 0.]
 [6. 0. 0. ... 0. 0. 0.]
 ...
 [8. 0. 0. ... 0. 0. 0.]
 [8. 0. 0. ... 0. 0. 0.]
 [7. 0. 0. ... 0. 0. 0.]]
```

```
(60000, 785)
```

```
In [10]: # Let's check the testing data
```

```
print(testData)
print('\n')
print(testData.shape)
```

```
[[0. 0. 0. ... 0. 0. 0.]
 [1. 0. 0. ... 0. 0. 0.]
 [2. 0. 0. ... 0. 0. 0.]
 ...
 [8. 0. 0. ... 0. 1. 0.]
 [8. 0. 1. ... 0. 0. 0.]
 [1. 0. 0. ... 0. 0. 0.]]
```

```
(10000, 785)
```

```
In [11]: # Let's visualize an image randomly using any random number from training d

# Importing library random to get any random number

import random
i = random.randint(1, trainData.shape[0])


# Displaying a random image using pixel values(reshaped)

plt.imshow(trainData[i, 1:].reshape(28, 28))
print(trainData[i, 0])

# Let's check all class values with labels for reference

# 0 = T-shirt/top
# 1 = Trouser
# 2 = Pullover
# 3 = Dress
# 4 = Coat
# 5 = Sandal
# 6 = Shirt
# 7 = Sneaker
# 8 = Bag
# 9 = Ankle boot
```

3.0



```
In [12]: # Let's view multiple images in a grid format


Grid_W = 10 # Number of columns
Grid_H = 10 # Number of rows

# Let's create subplot to display multiple images

fig, axes = plt.subplots(nrows=Grid_H, ncols=Grid_W, figsize = (18,18))
axes = axes.ravel()

# Let's loop through the training dataset to display random 100 images

for i in np.arange(0, Grid_W*Grid_H):
    index = np.random.randint(0, len(trainData))
    axes[i].imshow(trainData[index, 1:].reshape(28, 28))
    axes[i].set_title(trainData[index, 0])
    axes[i].axis('off')
```



Splitting and reshaping the data

```
In [13]: # Let's split the data into initial X_train and y_train data
# Normalizing each pixel value by dividing the same with 255
# i.e. Converting all pixel values between range 0 to 1

X_train_init = trainData[:, 1:] / 255
y_train_init = trainData[:, 0]

# Let's do the same for test data

X_test = testData[:, 1:] / 255
y_test = testData[:, 0]
```

```
In [14]: # Let's split the training data further to create training and validation data
        # Importing necessary library

        from sklearn.model_selection import train_test_split
        X_train, X_valid, y_train, y_valid = train_test_split(X_train_init, y_train_init,
                                                            test_size=0.2, random_state=42)
```

```
In [15]: X_train.shape
```

```
Out[15]: (48000, 784)
```

```
In [16]: X_valid.shape
```

```
Out[16]: (12000, 784)
```

```
In [17]: # Let's reshape the array similar to an image data e.g. (row, column, channel)

        X_train = X_train.reshape(X_train.shape[0], 28, 28, 1)
        X_test = X_test.reshape(X_test.shape[0], 28, 28, 1)
        X_valid = X_valid.reshape(X_valid.shape[0], 28, 28, 1)
```

```
In [18]: # Let's check the new shape of the training, testing and validation data

        X_train.shape
```

```
Out[18]: (48000, 28, 28, 1)
```

```
In [19]: X_test.shape
```

```
Out[19]: (10000, 28, 28, 1)
```

```
In [20]: X_valid.shape
```

```
Out[20]: (12000, 28, 28, 1)
```

Training a Neural Network Model

```
In [21]: import tensorflow as tf
```

First model

```
In [22]: #####
        # design your first model
        #####

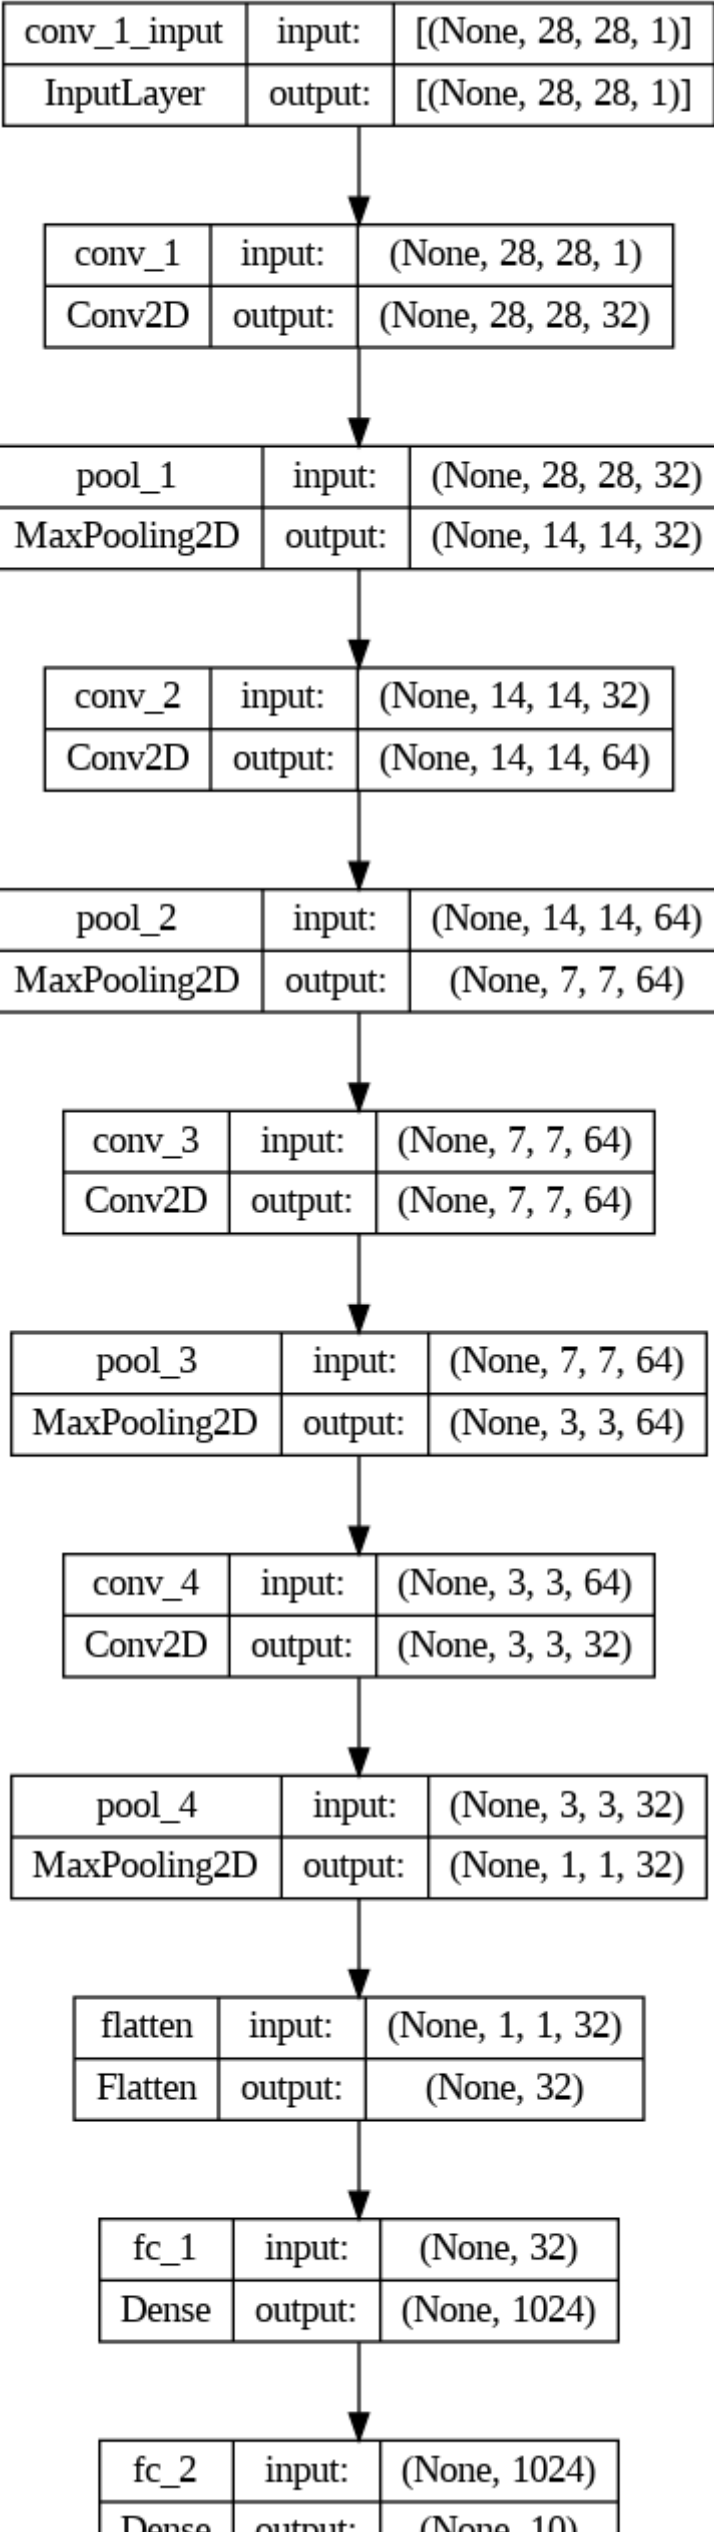
        model = tf.keras.Sequential()
        model.add(tf.keras.layers.Conv2D(filters=32, kernel_size=(5, 5), strides=(1, 1),
        model.add(tf.keras.layers.MaxPool2D(pool_size=(2, 2), name='pool_1'))
        model.add(tf.keras.layers.Conv2D(filters=64, kernel_size=(5, 5), strides=(1, 1),
        model.add(tf.keras.layers.MaxPool2D(pool_size=(2, 2), name='pool_2'))
        model.add(tf.keras.layers.Conv2D(filters=64, kernel_size=(5, 5), strides=(1, 1),
        model.add(tf.keras.layers.MaxPool2D(pool_size=(2, 2), name='pool_3'))
        model.add(tf.keras.layers.Conv2D(filters=32, kernel_size=(5, 5), strides=(1, 1),
        model.add(tf.keras.layers.MaxPool2D(pool_size=(2, 2), name='pool_4'))
        model.add(tf.keras.layers.Flatten())
        model.add(tf.keras.layers.Dense(units=1024, name='fc_1', activation='relu'))
        model.add(tf.keras.layers.Dense(units=10, name='fc_2', activation='softmax'))
        model.build(input_shape=(None, 28, 28, 1))
```

```
In [23]: model.compute_output_shape(input_shape=(10000, 28, 28, 1))
```

```
Out[23]: TensorShape([10000, 10])
```

```
In [24]: tf.keras.utils.plot_model(model, show_shapes=True, show_layer_names=True)
```

Out[24]:



Dense	Output.	(None, 10)
-------	---------	------------

```
In [25]: #####
# compile your first model
#####
model.compile(optimizer=tf.keras.optimizers.Adam(),loss=tf.keras.losses.Spa
```

```
In [26]: #####
# get summary of your model
#####
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv_1 (Conv2D)	(None, 28, 28, 32)	832
pool_1 (MaxPooling2D)	(None, 14, 14, 32)	0
conv_2 (Conv2D)	(None, 14, 14, 64)	51264
pool_2 (MaxPooling2D)	(None, 7, 7, 64)	0
conv_3 (Conv2D)	(None, 7, 7, 64)	102464
pool_3 (MaxPooling2D)	(None, 3, 3, 64)	0
conv_4 (Conv2D)	(None, 3, 3, 32)	51232
pool_4 (MaxPooling2D)	(None, 1, 1, 32)	0
flatten (Flatten)	(None, 32)	0
fc_1 (Dense)	(None, 1024)	33792
fc_2 (Dense)	(None, 10)	10250

=====
Total params: 249834 (975.91 KB)
Trainable params: 249834 (975.91 KB)
Non-trainable params: 0 (0.00 Byte)
=====

```
In [27]: #####
# fit your first model
#####
history = model.fit(X_train, y_train, epochs=3, validation_data=(X_valid, y
```

Epoch 1/3


/usr/local/lib/python3.10/dist-packages/keras/src/backend.py:5727: UserWarning: "`sparse_categorical_crossentropy` received `from_logits=True`, but the `output` argument was produced by a Softmax activation and thus does not represent logits. Was this intended?
output, from_logits = _get_logits(

1500/1500 [=====] - 173s 114ms/step - loss: 0.5109 - accuracy: 0.8086 - val_loss: 0.3320 - val_accuracy: 0.8799
Epoch 2/3
1500/1500 [=====] - 170s 113ms/step - loss: 0.3073 - accuracy: 0.8873 - val_loss: 0.2748 - val_accuracy: 0.8993
Epoch 3/3
1500/1500 [=====] - 168s 112ms/step - loss: 0.2586 - accuracy: 0.9038 - val_loss: 0.2840 - val_accuracy: 0.8963

```
In [28]: #####
# evaluate your first model and and plot your training curve and show class
#####
hist = history.history
x_arr = np.arange(len(hist['loss'])) + 1

fig = plt.figure(figsize=(12, 4))
ax = fig.add_subplot(1, 2, 1)
ax.plot(x_arr, hist['loss'], '-o', label='Train loss')
ax.plot(x_arr, hist['val_loss'], '--<', label='Validation loss')
ax.set_xlabel('Epoch', size=15)
ax.set_ylabel('Loss', size=15)
ax.legend(fontsize=15)
ax = fig.add_subplot(1, 2, 2)
ax.plot(x_arr, hist['accuracy'], '-o', label='Train acc.')
ax.plot(x_arr, hist['val_accuracy'], '--<', label='Validation acc.')
ax.legend(fontsize=15)
ax.set_xlabel('Epoch', size=15)
ax.set_ylabel('Accuracy', size=15)

plt.show()
```



```
In [29]: test_results = model.evaluate(X_test , y_test)
print(f"Test Accuracy is {test_results[1]*100}")
```

```
313/313 [=====] - 10s 30ms/step - loss: 0.2686 -
accuracy: 0.8985
Test Accuracy is 89.85000252723694
```

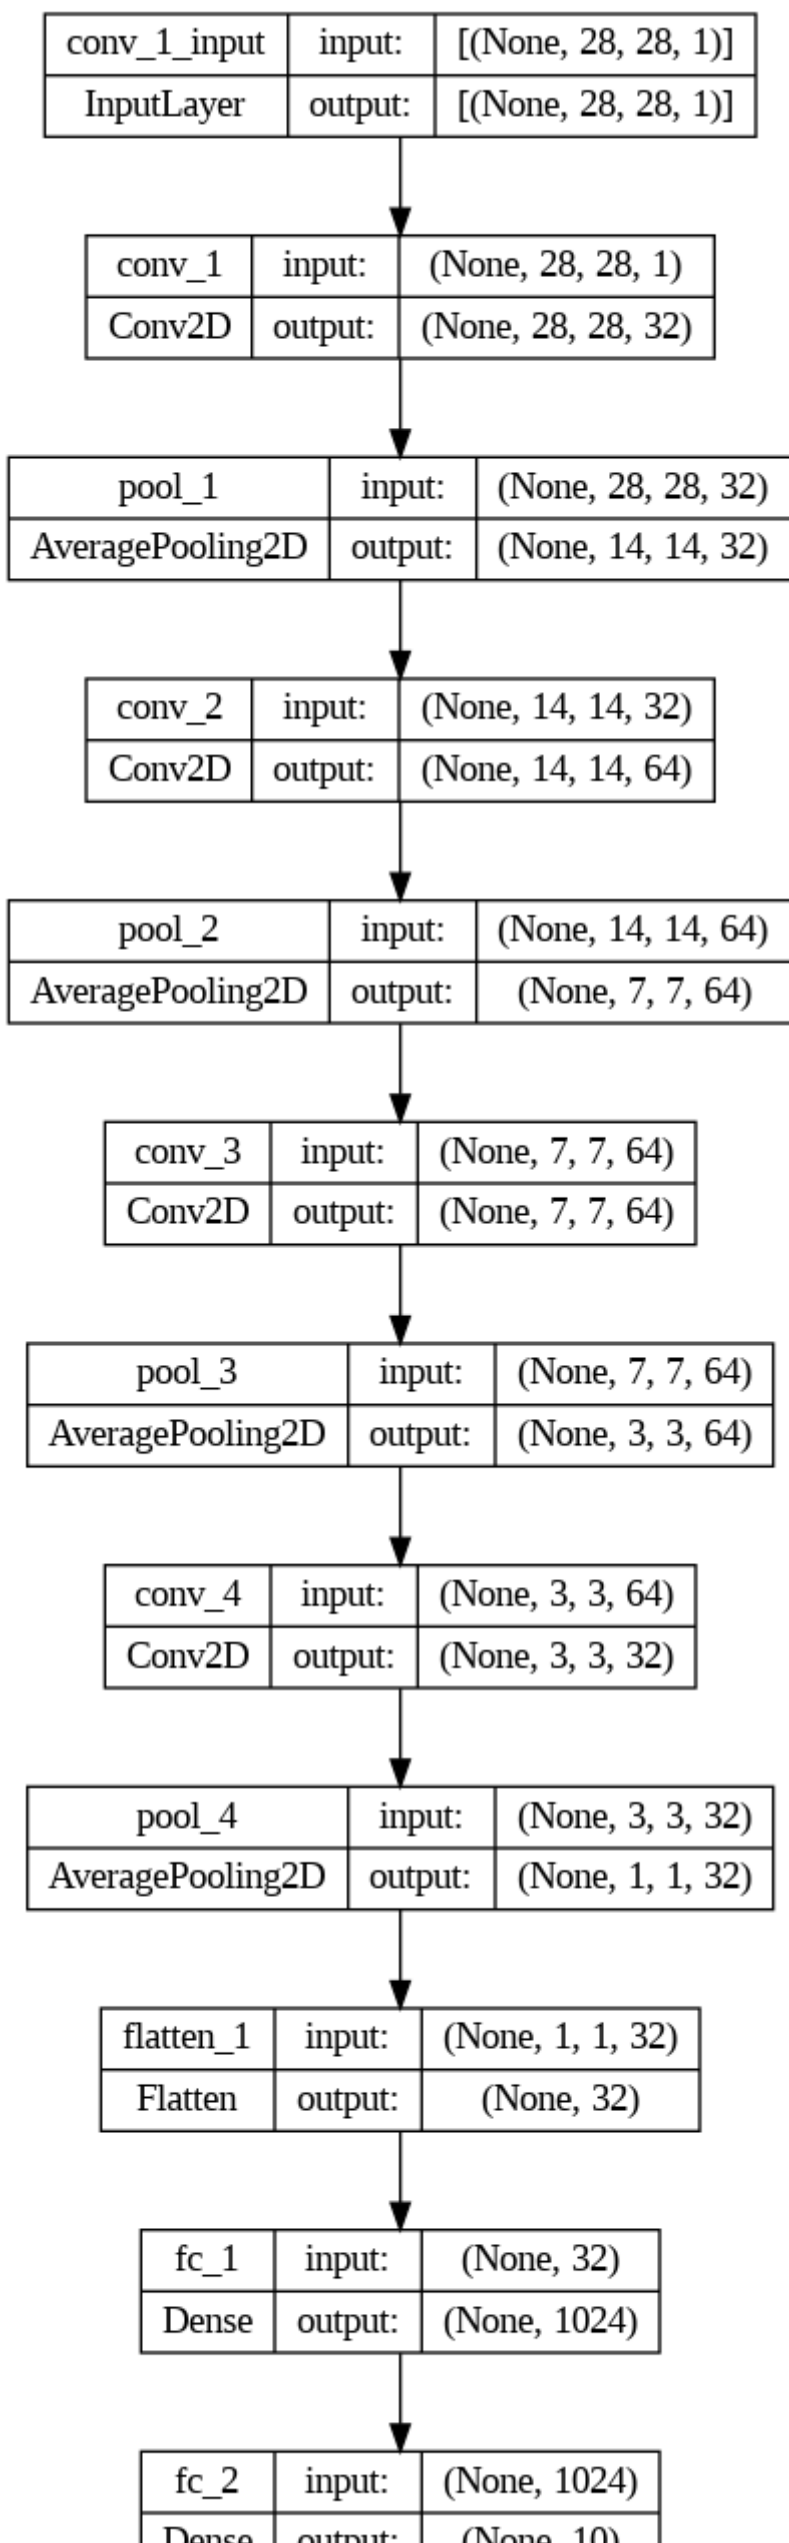
Second model

```
In [30]: #####
# design your second model
#####

model = tf.keras.Sequential()
model.add(tf.keras.layers.Conv2D(filters=32, kernel_size=(5, 5),strides=(1,
model.add(tf.keras.layers.AveragePooling2D(pool_size=(2, 2), name='pool_1')
model.add(tf.keras.layers.Conv2D(filters=64, kernel_size=(5, 5),strides=(1,
model.add(tf.keras.layers.AveragePooling2D(pool_size=(2, 2), name='pool_2')
model.add(tf.keras.layers.Conv2D(filters=64, kernel_size=(5, 5),strides=(1,
model.add(tf.keras.layers.AveragePooling2D(pool_size=(2, 2), name='pool_3')
model.add(tf.keras.layers.Conv2D(filters=32, kernel_size=(5, 5),strides=(1,
model.add(tf.keras.layers.AveragePooling2D(pool_size=(2, 2), name='pool_4')
model.add(tf.keras.layers.Flatten())
model.add(tf.keras.layers.Dense(units=1024, name='fc_1', activation='relu')
model.add(tf.keras.layers.Dense(units=10, name='fc_2',activation='softmax')
model.build(input_shape=(None, 28, 28, 1))
```

```
In [31]: tf.keras.utils.plot_model(model, show_shapes=True, show_layer_names=True)
```

Out[31]:



Dense	Output.	(None, 10)
-------	---------	------------

```
In [32]: #####
compile your first model
#####
del.compile(optimizer=tf.keras.optimizers.Adam(),loss=tf.keras.losses.Sparse
```

```
In [33]: #####
# get summary of your model
#####
model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv_1 (Conv2D)	(None, 28, 28, 32)	832
pool_1 (AveragePooling2D)	(None, 14, 14, 32)	0
conv_2 (Conv2D)	(None, 14, 14, 64)	51264
pool_2 (AveragePooling2D)	(None, 7, 7, 64)	0
conv_3 (Conv2D)	(None, 7, 7, 64)	102464
pool_3 (AveragePooling2D)	(None, 3, 3, 64)	0
conv_4 (Conv2D)	(None, 3, 3, 32)	51232
pool_4 (AveragePooling2D)	(None, 1, 1, 32)	0
flatten_1 (Flatten)	(None, 32)	0
fc_1 (Dense)	(None, 1024)	33792
fc_2 (Dense)	(None, 10)	10250

=====
Total params: 249834 (975.91 KB)
Trainable params: 249834 (975.91 KB)
Non-trainable params: 0 (0.00 Byte)
=====


```
In [34]: #####
# fit your second model
#####
history = model.fit(X_train, y_train, epochs=3, validation_data=(X_valid, y
```

Epoch 1/3
1500/1500 [=====] - 171s 113ms/step - loss: 0.6654 - accuracy: 0.7458 - val_loss: 0.4917 - val_accuracy: 0.8153
Epoch 2/3
1500/1500 [=====] - 159s 106ms/step - loss: 0.3990 - accuracy: 0.8522 - val_loss: 0.3612 - val_accuracy: 0.8673
Epoch 3/3
1500/1500 [=====] - 158s 105ms/step - loss: 0.3307 - accuracy: 0.8756 - val_loss: 0.2996 - val_accuracy: 0.8885

```
In [35]: #####
# evaluate your second model and and plot your training curve and show clas
#####
hist = history.history
x_arr = np.arange(len(hist['loss'])) + 1

fig = plt.figure(figsize=(12, 4))
ax = fig.add_subplot(1, 2, 1)
ax.plot(x_arr, hist['loss'], '-o', label='Train loss')
ax.plot(x_arr, hist['val_loss'], '--<', label='Validation loss')
ax.set_xlabel('Epoch', size=15)
ax.set_ylabel('Loss', size=15)
ax.legend(fontsize=15)
ax = fig.add_subplot(1, 2, 2)
ax.plot(x_arr, hist['accuracy'], '-o', label='Train acc.')
ax.plot(x_arr, hist['val_accuracy'], '--<', label='Validation acc.')
ax.legend(fontsize=15)
ax.set_xlabel('Epoch', size=15)
ax.set_ylabel('Accuracy', size=15)

plt.show()
```



```
In [36]: test_results = model.evaluate(X_test , y_test)
print(f"Test Accuracy is {test_results[1]*100}")

313/313 [=====] - 10s 30ms/step - loss: 0.2868 -
accuracy: 0.8899
Test Accuracy is 88.99000287055969
```

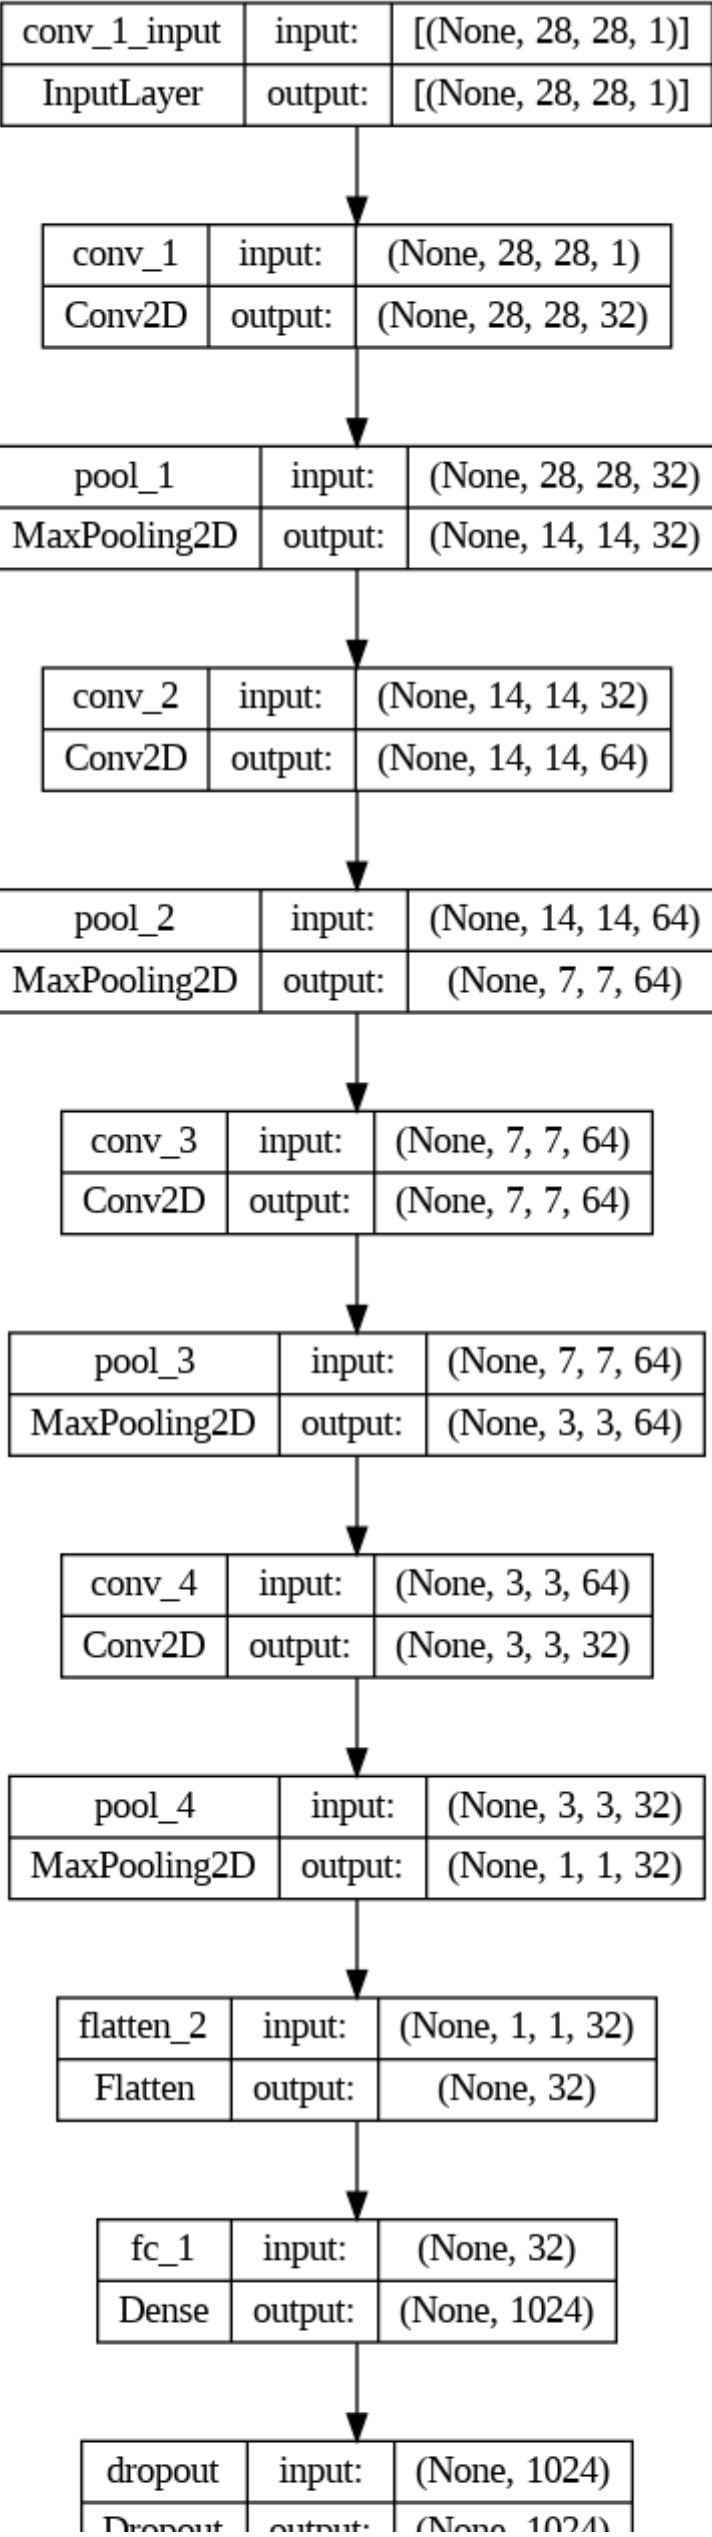
Third model


```
In [37]: #####
# design your third model
#####
model = tf.keras.Sequential()
model.add(tf.keras.layers.Conv2D(filters=32, kernel_size=(5, 5),strides=(1,
model.add(tf.keras.layers.MaxPool2D(pool_size=(2, 2), name='pool_1'))
model.add(tf.keras.layers.Conv2D(filters=64, kernel_size=(5, 5),strides=(1,
model.add(tf.keras.layers.MaxPool2D(pool_size=(2, 2), name='pool_2'))
model.add(tf.keras.layers.Conv2D(filters=64, kernel_size=(5, 5),strides=(1,
model.add(tf.keras.layers.MaxPool2D(pool_size=(2, 2), name='pool_3'))
model.add(tf.keras.layers.Conv2D(filters=32, kernel_size=(5, 5),strides=(1,
model.add(tf.keras.layers.MaxPool2D(pool_size=(2, 2), name='pool_4'))
model.add(tf.keras.layers.Flatten())
model.add(tf.keras.layers.Dense(units=1024, name='fc_1', activation='relu'))
model.add(tf.keras.layers.Dropout(rate = 0.5))
model.add(tf.keras.layers.Dense(units=10, name='fc_2',activation='softmax'))
model.build(input_shape=(None, 28, 28, 1))
```



```
In [38]: tf.keras.utils.plot_model(model, show_shapes=True, show_layer_names=True)
```

Out[38]:





```
In [39]: #####
# compile your first model
#####
model.compile(optimizer=tf.keras.optimizers.Adam(), loss=tf.keras.losses.Spa
```

```
In [40]: #####
# get summary of your model
#####
model.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
=====		
conv_1 (Conv2D)	(None, 28, 28, 32)	832
pool_1 (MaxPooling2D)	(None, 14, 14, 32)	0
conv_2 (Conv2D)	(None, 14, 14, 64)	51264
pool_2 (MaxPooling2D)	(None, 7, 7, 64)	0
conv_3 (Conv2D)	(None, 7, 7, 64)	102464
pool_3 (MaxPooling2D)	(None, 3, 3, 64)	0
conv_4 (Conv2D)	(None, 3, 3, 32)	51232
pool_4 (MaxPooling2D)	(None, 1, 1, 32)	0
flatten_2 (Flatten)	(None, 32)	0
fc_1 (Dense)	(None, 1024)	33792
dropout (Dropout)	(None, 1024)	0
fc_2 (Dense)	(None, 10)	10250
=====		
Total params: 249834 (975.91 KB)		
Trainable params: 249834 (975.91 KB)		
Non-trainable params: 0 (0.00 Byte)		

```
In [41]: #####
# fit your third model
#####
history = model.fit(X_train, y_train, epochs=3, validation_data=(X_valid, y

Epoch 1/3


/usr/local/lib/python3.10/dist-packages/keras/src/backend.py:5727: UserWarning: "`sparse_categorical_crossentropy` received `from_logits=True`, but the `output` argument was produced by a Softmax activation and thus does not represent logits. Was this intended?
  output, from_logits = _get_logits(

1500/1500 [=====] - 170s 113ms/step - loss: 0.5178 - accuracy: 0.8063 - val_loss: 0.3708 - val_accuracy: 0.8738
Epoch 2/3
1500/1500 [=====] - 167s 111ms/step - loss: 0.3118 - accuracy: 0.8857 - val_loss: 0.2780 - val_accuracy: 0.8996
Epoch 3/3
1500/1500 [=====] - 177s 118ms/step - loss: 0.2636 - accuracy: 0.9037 - val_loss: 0.2836 - val_accuracy: 0.8957
```

```
In [42]: #####
# evaluate your third model and and plot your training curve and show class
#####
hist = history.history
x_arr = np.arange(len(hist['loss'])) + 1

fig = plt.figure(figsize=(12, 4))
ax = fig.add_subplot(1, 2, 1)
ax.plot(x_arr, hist['loss'], '-o', label='Train loss')
ax.plot(x_arr, hist['val_loss'], '--<', label='Validation loss')
ax.set_xlabel('Epoch', size=15)
ax.set_ylabel('Loss', size=15)
ax.legend(fontsize=15)
ax = fig.add_subplot(1, 2, 2)
ax.plot(x_arr, hist['accuracy'], '-o', label='Train acc.')
ax.plot(x_arr, hist['val_accuracy'], '--<', label='Validation acc.')
ax.legend(fontsize=15)
ax.set_xlabel('Epoch', size=15)
ax.set_ylabel('Accuracy', size=15)

plt.show()
```



```
In [43]: test_results = model.evaluate(X_test , y_test)
print(f"Test Accuracy is {test_results[1]*100}")

313/313 [=====] - 8s 26ms/step - loss: 0.2779 - accuracy: 0.8973
Test Accuracy is 89.73000049591064
```